

CTA Strategy · Course Notes

Harry Chen

Contents

1 · What Is a CTA Strategy	2
1.1 What a CTA Trades	2
1.1.1 Index, ETF, or future	2
1.1.2 Why futures rather than ETFs	4
1.2 Momentum, the Source of the Return	4
1.2.1 What momentum is	4
1.2.2 Why momentum might work	4
1.3 Long and Short, the Two Sides of a Position	5
1.3.1 What the two sides mean	5
1.3.2 How a short is borrowed	5
1.3.3 Why the short leg is the fragile one	5
1.3.4 How a short is represented	6
1.4 Who Else Is Trading	7
Background	7
Appendix · Notation	8
5 · Modelling	10
5.1 Why one split is not a backtest	10
5.1.1 The textbook split	10
5.1.2 What the single split costs	10
5.1.3 The fix that is not one	11
5.2 The walk-forward ladder	12
5.2.1 How long a segment should be	12
5.2.2 The slide	12
5.2.3 One rung, in full	12
5.3 Two validation layers	13
5.3.1 Why the inner layer cannot be skipped	14
5.3.2 What only the outer block may decide	14
5.3.3 The optimizer's parameters	15
5.4 Preparing a signal for a model	15
5.4.1 Clipping: keep the observation, cap the value	15

5.4.2	Standardizing: put every signal on one scale	17
5.4.3	Which mean and standard deviation you may use	18
Appendix	Notation	18

1 · What Is a CTA Strategy

1.1 What a CTA Trades

Definition 1.1 (CTA strategy). A *CTA* (*Commodity Trading Advisor*) strategy is a rule-based strategy that trades **futures** systematically across asset classes—equities, fixed income, commodities, currencies—targeting **absolute return**, independent of market direction.

Note. The name is a historical artifact. Modern CTAs are not commodity-only; typical markets are equity-index futures, Treasuries, FX, energy, metals, and agriculture.

1.1.1 Index, ETF, or future

“S&P 500” names three different objects. Only two of them can be traded.

Layer	Example	What it is	Tradeable
Index	SPX, NDX	A published number computed from its constituents	No
ETF	SPY, VOO, QQQ	A fund holding the constituents; its own shares trade	Yes
Future	ES, NQ	A contract to settle at a future date; holds nothing	Yes

Definition 1.2 (Index).

$$I_t = \frac{M_t}{D_t}, \quad M_t = \sum_i N_i P_{it}$$

where P_{it} is the price of constituent i , N_i its float-adjusted share count, M_t the total float capitalisation, and $D_t > 0$ the *divisor*.

Claim 1.3. Index returns are the capitalisation-weighted average of constituent returns; the divisor sets only the level.

Proof. Both halves follow from I_t being a *sum divided by a constant*. A constituent’s dollar move is N_i times its price change, so its share of the index move is its share of total capitalisation:

$$r_I = \sum_i w_{it} r_i, \quad w_{it} = \frac{N_i P_{it}}{M_t}, \quad \sum_i w_{it} = 1.$$

And D_t divides I_t at both dates, so it cancels from the ratio—it can shift the level anywhere without touching a single return. $\square$

Note (Two consequences). **The top names dominate.** At $w \approx 0.07$ against $w \approx 0.0002$, the largest constituent carries roughly **350 times** the index impact of the smallest; this is not an average of 500 *prices*. **And the level is arbitrary.** D_t cancels from returns, so the scale is whatever the base period set it to—for the S&P 500, $1941-43 = 10$. **Compare returns, never levels.**

A price level is meaningful only against that same series' own history. Across tickers it carries no information at all—what each ETF's share price happens to be was fixed by how the fund was structured at inception. Dividing every series by its own first observation removes that arbitrariness:

Ticker	Raw close, start	Raw close, end	Rebased	Return
SPY	325.05	740.83	227.9	+127.9%
TLT	137.10	87.35	63.7	−36.3%
GLD	143.97	368.49	255.9	+155.9%

SPY has the highest closing price on every single day of the sample, yet GLD returned more. The end level alone cannot tell you this; you need the start it is measured against.




Figure 1.1: The same three ETFs, twice. Left, raw closing prices: SPY runs highest for the entire sample. Right, each series divided by its own first value and multiplied by 100: GLD now finishes highest and TLT ends below the line it started on. The ordering of the levels and the ordering of the returns are different orderings.

1.1.2 Why futures rather than ETFs

Reason	Futures	ETFs
Symmetry	Shorting costs nothing extra	Shorting needs borrowed shares, a fee, availability
Coverage	Deep markets in all four asset classes	Poor in commodities and FX
Capital efficiency	Margin is roughly 5 to 10 percent of notional; spare cash earns collateral yield	Full notional tied up

Symmetry is decisive: a strategy that must go short as freely as long cannot tolerate the asymmetry, and §1.3.3 shows why the short leg is the fragile one. The cost is that futures expire, so a continuous series has to be stitched together from contracts—see 100 · The Dataset.

1.2 Momentum, the Source of the Return

1.2.1 What momentum is

Definition 1.4 (Momentum). *Momentum* is the empirical tendency for recent relative performance to persist: assets that outperformed over the past one to twelve months tend to keep outperforming, and recent laggards tend to keep lagging.

Note. The claim is about *returns*, not price levels—it concerns the ordering of performance across assets and across time. That is why one momentum rule can be applied to markets with nothing else in common, from Treasury futures to natural gas.

Turning this into a number—a lookback window, a normalization, a test that it carries information—is 2 · Testing a Signal’s subject. Here it is only the premise.

1.2.2 Why momentum might work

Two contested explanations, no proof. Momentum has been durably profitable regardless.

- **Information diffusion.** Unglamorous macro news spreads gradually, creating sustained pressure. Possibly weakened since 2008, as computing and data distribution got faster.
- **Selective momentum capture.** Winners keep winning and losers keep losing, because large positions take time to build and unwind, investors chase winners, and major events do not resolve in a single day.

2 · Testing a Signal turns this from a story into something testable.

1.3 Long and Short, the Two Sides of a Position

1.3.1 What the two sides mean

Definition 1.5 (Long). Buy first, sell later. Profit is the sell price minus the buy price.

Definition 1.6 (Short). Borrow and sell first, buy back and return later. Profit is the sell price minus the buy-back price.

Note. Running both deploys capital fully—a 150/50 or 200/100 book rather than idle cash—and lets the strategy express a negative view, which a long-only book cannot.

1.3.2 How a short is borrowed

A short is not selling out of thin air—the shares are borrowed before they are sold:

1. **Borrow** N shares from a lender—a broker, or a long-term holder.
2. **Sell** them now, and receive cash.
3. **Buy back** N shares later, from the market.
4. **Return** the N shares, closing the position.

Note. You never create shares. They are fungible, so returning “ N shares of SPY”—not specific certificates—settles the debt. **Lenders are paid.** Long-term holders earn a **lending fee** on stock they were going to sit on anyway. **And the broker sits in the middle**, matching the two sides and holding margin; the market for this is **securities lending**.

→ The full mechanics beneath this—who holds the shares mid-loan, who receives the dividend, hard-to-borrow rates, and how a short squeeze runs away: §10 of the `market-101-foundations` note in `market_knowledge/`.

1.3.3 Why the short leg is the fragile one

Claim 1.7. For a static position, a long’s loss is bounded by its outlay; a short’s is not.

Proof. With signed share count s , entry price $P_0 > 0$ and exit price P_T , the payoff is

$$\text{PnL}(P_T) = s(P_T - P_0), \quad P_T \in [0, \infty)$$

the domain being half-open because a price cannot go negative.

For a **long**, $s > 0$: PnL is increasing in P_T , so its infimum is at the left endpoint, $\text{PnL} \geq s(0 - P_0) = -sP_0$ —exactly the amount paid. For a **short**, $s < 0$: PnL is decreasing in P_T , and the domain has no right endpoint, so $\text{PnL} \rightarrow -\infty$ as $P_T \rightarrow \infty$. □

The asymmetry is in the **domain**, not the payoff—both are lines of slope $|s|$. The floor exists only because prices are floored at zero.




Figure 1.2: Profit and loss per dollar of exposure against the terminal price P_T . The long and short payoffs are lines of equal and opposite slope crossing at the entry price P_0 ; the long's line stops at -100 percent where the price reaches zero, while the short's continues downward without limit as the price rises. Illustrative.

Note. Hence margin requirements and forced buy-ins—and why a 150/50 book, 200 percent gross, is **two times leverage** rather than free money.

1.3.4 How a short is represented

A short needs no separate machinery. Carry the position as a **signed quantity** and let the sign do the work: a negative share count valued at the market price is a **negative position value**, which rises exactly when the price falls.

Note (Sign convention). A negative position value is a **liability**—a debt owed—not a loss already taken.

Note (Static versus constant-exposure shorts). The proof above holds the share count fixed at entry. A book that instead maintains a constant *dollar* exposure buys the short back down as it moves against it, so the unbounded loss is never realised. That daily rebalancing is an **implicit risk control**: the proof is right, but a constant-exposure book is not running the position it describes.

→ How this project implements it, with the accounting derivation and the measured gap between the two: `Backtests.md` in `Backtest_prototype/`.

1.4 Who Else Is Trading

Participant	Examples	Typical hold
Index / passive funds	Vanguard, BlackRock	Years to decades
Active managers	Fidelity, PIMCO	Monthly to quarterly
Hedge funds	—	Intraday to several days
Market makers	Optiver, Citadel Securities, IMC, SIG	Seconds to minutes
Noise traders	Retail, uninformed flow	No consistent horizon

Background

General finance knowledge needed to read this chapter, not part of its argument.

What is the relationship between futures/options and asset classes?

They answer different questions, so they are not alternatives to one another.

- **Asset class**—*what* underlying market you are exposed to.
- **Instrument**—*how* you hold that exposure: cash, futures, or options.

Futures are not an asset class, and options are not a rival one. Every asset class trades in every instrument:

Asset class	Typical underlying	Futures	Options
Equities	S&P 500, Nasdaq 100	ES, NQ	options on ES
Fixed income	US Treasuries, Bunds	ZN (10-year), ZB (30-year)	options on ZN
Commodities	Crude oil, gold, corn	CL, GC, ZC	options on CL
Currencies (FX)	EUR/USD, USD/JPY	6E, 6J	options on 6E

The same exposure comes in three wrappers: equity exposure is bought in cash as SPY, in futures as ES, or in options as an ES call. So in the opening definition—“trades **futures** systematically **across asset classes**”—the three parts are independent: **futures** is the wrapper, **across asset classes** is the breadth of underlying markets, and **systematically** means by rule or model rather than discretion.

Which asset class does a bond belong to?

Fixed income. Treasuries, corporate, municipal and high-yield bonds all sit there.

- “Fixed” describes the **contract**, meaning a scheduled coupon and the return of principal—not a fixed price.
- Bond prices still move, with interest rates, credit risk and time to maturity. That variation is what makes them a trend market at all.
- Treasury futures trade at 2-, 5-, 10- and 30-year maturities, so a CTA picks the *maturity* it wants exposure to.

Why do CTAs use futures rather than options?

	Futures	Options
Exposure	Linear —profit and loss move one-for-one with the underlying	Non-linear; depends on strike and moneyness
Extra drivers	None	Implied volatility, time decay, strike
What you are betting on	Direction	Direction and volatility

An option position is partly a volatility trade whether you intended it or not. Linear exposure is what lets one risk framework size 37 positions on the same scale.

Example 1.8. A trend follower might simultaneously hold long ES, short ZN, long GC, short 6J—all four are futures, and all four are different asset classes.

Appendix · Notation

Throughout, t is the date and i indexes an index’s constituents.

Symbol	Means	First used
I_t	the index level on date t —a published number, not a tradeable object	§1.1.1
P_{it}, N_i	constituent i ’s price on date t , and its float-adjusted share count	§1.1.1
M_t, D_t	the total float capitalisation, and the divisor that turns it into the printed level	§1.1.1
w_{it}	constituent i ’s share of total capitalisation—its weight in the index return	§1.1.1

Symbol	Means	First used
r_I, r_i	the index's return and constituent i 's return, over the same period	§1.1.1
s	the signed share count of a position: positive is long, negative is short	§1.3.3
P_0, P_T	the entry and exit price of that position	§1.3.3
$\text{PnL}(P_T)$	the position's profit and loss as a function of where the price ends up	§1.3.3
N	the number of shares borrowed to open a short	§1.3.2

Note (Collisions to watch). s here is a **signed share count**, not 2's s , which indexes the asset. w_{it} is a **constituent's weight inside an index**, fixed by capitalisation and not chosen by anyone; 2's $w_{s,t}$ and 5 · §3.3's w_i are **portfolio** weights, which are decisions. r_i is a constituent's return, whereas elsewhere in the course r carries an asset and a date subscript and means the forward return being predicted.

5 · Modelling

5.1 Why one split is not a backtest

5.1.1 The textbook split

Definition 5.1 (Split). A *split* cuts the sample by date into three contiguous blocks: **train**, on which coefficients are estimated; **validation**, on which choices between fitted models are made; and **test**, which is scored once and never chosen on.




Figure 5.1: One split against the walk-forward ladder. Above, a single timeline cut once into train, validation and test. Below, the same history divided into twelve equal segments, with the train–validation–test frame starting one segment later on each row—so each row’s test segment is the next row’s validation segment.

The order matters and is not a convention: the blocks run forward in time, because a model that was fitted on next year and scored on last year has learned something no live book could have known.

5.1.2 What the single split costs

Three defects, and they compound.

Defect	What goes wrong
One fit discards time	A single set of coefficients over several years asserts that those years were one market. Early in the pandemic the world locked down at once and nobody wanted oil; the consensus was that demand had permanently moved online. It had not—the economy reopened, and energy then flew on war. A coefficient estimated across both states describes neither.
Almost nothing is left to score	Give validation a year and test a year and validation is roughly 250 dates. A year is also too little market: the regimes it contains are whichever ones happened to fall inside it, so the split never asks what the strategy does when the regime turns. A maximum drawdown measured over one calm year is not a risk estimate.
The sample is small and the signal is faint	A school dataset with a thousand points and an R^2 of 0.5 is ordinary, and in physics 0.9 is ordinary. In finance 0.1 to 0.2 is already a good R^2 —see 9. Weak true dependence plus few observations is precisely the regime in which a fit lands on noise: a line through ten points drawn from a cloud reaches R^2 0.8 easily.

5.1.3 The fix that is not one

Claim 5.2. Duplicating the sample changes no coefficient, and inflates every t-statistic by roughly $\sqrt{2}$.

Proof. Write the fit as $b = (X^T X)^{-1} X^T y$, with X the $n \times p$ matrix of signal values, y the forward returns, and b the estimated coefficients. Stacking the sample twice gives $X_2^T X_2 = 2 X^T X$ and $X_2^T y_2 = 2 X^T y$, so

$$b_2 = (2 X^T X)^{-1} 2 X^T y = b.$$

The residuals repeat as well, so the residual sum of squares doubles while the degrees of freedom go from $n - p$ to $2n - p$; for n large the variance estimate s^2 is therefore essentially unchanged. But the reported covariance is $s^2(X_2^T X_2)^{-1} = s^2(X^T X)^{-1}/2$ —halved. Standard errors fall by $\sqrt{2}$ and t-statistics rise by the same factor, on evidence that never arrived. $\square$

Note. This is the standard interview question, and the reason it is asked is that people reach for it. Manufacturing rows is overfitting with extra steps; the answer to too few observations is to score more of the ones you have, which is §5.2.

5.2 The walk-forward ladder

5.2.1 How long a segment should be

Definition 5.3 (Segment). A *segment* is a short contiguous block of dates—here **two to three months**—treated as one regime.

The length is set by how fast the market can change its mind, and the market is made of people. Within a day a regime shift is implausible: individuals change their view constantly, the crowd does not. Within a month, after a run of large news, one shift is plausible. Within a year or two there is room for several, because that is enough time for the same participants to change their minds repeatedly. So a segment must be short enough that one segment is approximately one regime, and long enough to hold a usable sample.

5.2.2 The slide

Cut the history into segments 1, 2, 3, ... and slide a frame across them:

Iteration	Train	Validation	Test
1	1–3	4	5
2	2–4	5	6
3	3–5	6	7

Claim 5.4. The ladder repairs all three defects of §5.1.2 at once.

Each fit now sees one regime’s worth of data and is re-estimated as the market moves on, so time is no longer averaged away. Every segment past the first frame is eventually scored out of sample, so the out-of-sample record grows with the length of the history instead of being fixed at one year. And a regime shift that lands inside a validation segment is now a *test* the strategy has to pass, rather than an event the single split happened to miss.

Note. Read the table down its two right columns: **each iteration’s test segment is the next iteration’s validation segment.** Nothing is discarded between rungs, and no segment is ever scored twice under the same coefficients.

5.2.3 One rung, in full

A rung is not a single fit. Train ends, a validation segment intervenes, and only then does test begin—so a model fitted on train alone is asked to predict across a gap of market it has never seen, and pays for it.

The repair is the two-step shape that Lasso already imposes. Lasso adds an L_1 penalty to the least-squares objective,

$$\min_{\beta} \sum_t \left(y_t - \sum_j \beta_j x_{j,t} \right)^2 + \lambda \sum_j |\beta_j|$$

where y_t is the forward return on date t , $x_{j,t}$ signal j 's value on that date, β_j its coefficient, and $\lambda \geq 0$ the penalty strength. The penalty drives coefficients exactly to zero, which makes it a **feature selector**; and because the surviving coefficients are biased toward zero by that same penalty, the standard practice is to **refit** the survivors without the penalty before predicting.




Figure 5.2: One iteration, in four stages: fit on train; select features on the validation segment; refit the survivors on a window shifted forward so that it ends where the prediction starts; apply once to test, beyond the dashed prediction date.

So each rung runs: **fit on train**, **select features on the validation segment**, **refit the survivors on a window ending where the prediction starts**, **apply once to test**.

Note (Two house styles). The refit window is either the training window shifted forward, or the training window plus the validation segment. The sample sizes differ little and both are used in industry; it is a preference, not a correctness question. What is not optional is that the window ends at the prediction date—see §5.4.3.

5.3 Two validation layers

The validation *segment* inside each rung and the large held-out validation *block* at the end of the history are different objects with different powers.




Figure 5.3: The two layers. The inner layer is the miniature ladder inside the training block, which chooses which features survive; the outer layer is the ringed held-out validation block, which chooses the model class, its hyperparameters, and the optimizer's parameters.

Layer	May choose	May not
Inner —the validation segment inside each rung	which features survive, and the refit that follows	anything held constant across rungs
Outer —the held-out validation block	the model class , its hyperparameters (λ), and the optimizer's parameters	coefficients—those are fitted on train
Test	nothing	everything

5.3.1 Why the inner layer cannot be skipped

Exploration already looked at which signals worked *in the training period*. Of course the fit then describes that period well: momentum was selected partly because it was seen to work there. That says nothing about whether it was useful going in, and the only way to find out is to select on dates the selection did not see—which is what the validation segment is for.

Note. A signal's usefulness is not a constant either. When more signals enter the book, or when one starts to fail, momentum should no longer carry the weight it carried alone. Because the ladder re-selects at every rung, that adjustment happens on its own rather than being asserted once.

5.3.2 What only the outer block may decide

λ is not fitted, it is searched: run the ladder once per value on a grid—0.1, 0.5, 0.7, 1—and compare. The same is true of the model class itself (Lasso, Ridge, a decision tree). Both comparisons need dates that no rung of the ladder consumed, which is exactly what the held-out block is.

5.3.3 The optimizer’s parameters

The prediction is not yet a portfolio. Weights come from a mean–variance problem: choose w to

$$\max_w E[r_p] - \gamma \cdot \text{Var}[r_p]$$

subject to constraints such as

$$\sum_i w_i = 1.0 \quad \sum_i |w_i| \leq 2.0 \quad |w_i| \leq 0.20$$

—net exposure 100 percent, gross exposure at most 200 percent, and no single asset above 20 percent—where w_i is the weight on asset i , $E[r_p]$ and $\text{Var}[r_p]$ the portfolio’s expected return and variance, and $\gamma \geq 0$ the **risk-aversion coefficient**.

Two of these are free parameters, and neither can be reasoned to:

- γ . Nobody can say from first principles whether a risk appetite is 0.3, 0.5 or 1. The number has no interpretable units—it is a trade rate between two quantities measured in different things—so it is chosen by running the grid and reading the result.
- **The per-asset cap.** With 37 assets, a 20 percent cap may be too tight for the book to hold enough of whichever names are paying; 25 or 30 percent may be the right ceiling. Gross and net exposure, by contrast, are set by mandate and are not tuned.

Note. All of this tuning belongs in the **outer** block. Tuning γ inside the rungs would let a portfolio parameter be chosen on the same dates the features were chosen on, and the held-out block would then be measuring a strategy that had already been fitted to it.

5.4 Preparing a signal for a model

Two transformations stand between a raw signal and a model, and both exist to stop the fit from tracking things that are not the market.

5.4.1 Clipping: keep the observation, cap the value

Claim 5.5. A ranked signal does not fill its range evenly—the mass piles at both ends—and those ends move the fit far more than their count suggests.

The first half is empirical: momentum in particular pins at an extreme whenever a name runs, so the histogram is a barbell rather than a plateau. The second half is mechanical. In a least-squares fit each observation’s pull on the slope grows with its distance from the mean of x , so a handful of far points can swing a coefficient that




Figure 5.4: Left, the share of observations by decile of the ranked signal—high at the first and last deciles, low in the middle, against the dashed line at 0.1 that an even spread would give. Right, a signal series against a shaded rolling percentile band, with the spikes that reach outside it pulled back onto its edge.

hundreds of central points barely move. The coefficient then jumps for reasons that have nothing to do with the market.

Deleting the far points, which is the textbook move, fails twice here:

- **The sample is already small.** Every deletion sharpens §5.1.2’s third defect, and in a messy sample it is not even clear which points are outliers.
- **The tail carries information.** What the book does when a signal goes extreme is exactly what the model should learn. Delete those dates and it never can.

Definition 5.6 (Winsorization). Replacing any value outside a chosen quantile band by the nearer edge of the band:

$$x_{j,t}^{\text{clip}} = \min \left(\max \left(x_{j,t}, Q_{j,t}^{\text{lo}} \right), Q_{j,t}^{\text{hi}} \right)$$

where $Q_{j,t}^{\text{lo}}$ and $Q_{j,t}^{\text{hi}}$ are quantiles of signal j over the trailing L dates—1 and 99 percent, or 2.5 and 97.5 percent; there is no canonical pair.

Note (Compare raw to raw). The band for date t must be estimated from the **un-clipped** history. Feed yesterday’s capped value into today’s quantile and the band walks inward on itself, a little further every step. Keep the raw column and produce the clipped one beside it.

Note (Why the band rolls). A value that is extreme against the last three months may be unremarkable six months later, once the market has moved to it. A fixed threshold pins such a signal at the 99th percentile forever; a rolling one releases it as soon as the surrounding distribution catches up.

→ `rolling`, `quantile` and `clip` as pandas calls: 8 · Toolbox.

5.4.2 Standardizing: put every signal on one scale

Momentum lives in roughly $[-0.05, 0.05]$. VIX lives in $[0, 100]$. Fit both and the coefficients absorb the mismatch.

Claim 5.7. Coefficient magnitudes carry no information about a signal’s importance until the signals share a scale.

Proof. Replace x_j by $c x_j$ for some $c > 0$. The fitted values are unchanged if β_j is replaced by β_j/c , and least squares chooses fitted values—so the fit is identical and the coefficient is whatever the units make it. A comparison of β magnitudes across differently scaled signals is therefore a comparison of units. $\square$

The consequence is not merely aesthetic. Every input is an **estimate** and carries error: a momentum measured from returns is the residue of participants entering and leaving, and a jump in it is as likely to be noise as news. A large β multiplies that noise into the prediction, while the small β on the wide-ranging signal cannot respond when it genuinely moves.

Definition 5.8 (Standardized signal).

$$z_{j,t} = \frac{x_{j,t} - \mu_j}{\sigma_j}$$

with μ_j and σ_j the mean and standard deviation of signal j over the estimation window. Dividing by a rolling standard deviation, or by a cross-sectional sum, are alternatives; what matters is that **every signal is treated the same way**.




Figure 5.5: Three coefficient paths over a year of refits: momentum decays from high toward zero, a volatility-regime signal rises from zero to become the largest, and MACD stays between them. The shaded stretches mark where momentum carries the book and where the lean rotates to the regime signal.

Three things follow, in increasing order of usefulness:

Consequence	Why
Coefficients become comparable	With units removed, β_j measures response per standard deviation of signal—the same question for every signal.
The coefficient path becomes a regime read	Plot every β_j from every rung on one axis. Momentum dominating one stretch and volatility another is the change in market sentiment, in a form you can look at.
Errors partly cancel	On one scale, with comparable coefficients and errors that are close to independent across signals, the noise in the combined prediction averages down as signals are added, instead of being dominated by whichever signal has the largest units.

5.4.3 Which mean and standard deviation you may use

μ_j , σ_j and the quantile band all have to be estimated from somewhere, and that choice is where look-ahead enters.

The rule follows from **where the model is built**: at the *end* of the training window. Every date inside that window is history at that moment, so standardizing with the whole training window's mean and standard deviation is legitimate—even though, from the point of view of an individual date early in the window, the statistic used to scale it was computed partly from its own future.

At this point of construction	These statistics are available
End of the training window	the whole training window
End of the training window	never the validation segment, and never test
End of a refit window	the whole refit sample, validation segment included

Note. The distinction is between *a signal's* information set and *the model's*. A signal computed on date t may use only data up to t , because it has to be computable live. The model is not computed on date t ; it is computed once, at the end of its window, and then applied forward. Confusing the two costs either a leak or an estimate worse than the one you were entitled to.

Note. This is the failure worth being paranoid about. In a time series the smallest leak makes the result beautiful, and the gap between a beautiful backtest and the live book is where the money goes.

Appendix · Notation

Throughout, t is the date, i indexes assets and j indexes signals.

Symbol	Means	First used
$x_{j,t}, x_{j,t}^{\text{clip}}, z_{j,t}$	signal j on date t : raw, after clipping, after standardizing	§5.2.3, §5.4.1, §5.4.2
y_t, β_j	the forward return being predicted, and the coefficient the fit puts on signal j	§5.2.3
X, y, b, n, p, s^2	the design matrix, the response, the estimated coefficients, the number of dates, the number of signals, the residual variance estimate	§5.1.3
λ	the L_1 penalty strength—searched on a grid, not fitted	§5.2.3
L	the trailing window, in dates, that the clip band is estimated over	§5.4.1
$Q_{j,t}^{\text{lo}}, Q_{j,t}^{\text{hi}}$	the low and high quantiles of that window	§5.4.1
μ_j, σ_j	the mean and standard deviation used to standardize signal j	§5.4.2
$w_i, \gamma, E[r_p]$	the portfolio weight on asset i , the risk-aversion coefficient, and the portfolio's expected return	§5.3.3

Note (Collisions to watch). σ_j here is one **signal's** dispersion, used only to rescale it—not 2 · §4's $\sigma_{s,t}$, an asset's trailing volatility, and not 4 · §1.1's σ , the amplitude of the tape. β_j is a coefficient on a **signal**, as in 4 · §3.2, never a market beta. λ is the penalty and γ the risk aversion: the lecture wrote both as λ and then renamed the second, and they are unrelated. L is a trailing window as in 4 · §2.1, here over a signal rather than over returns.